

# classic-pred

May 31, 2025

```
[1]: import pandas as pd
import numpy as np
import os
import warnings
from sklearn.model_selection import KFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
import matplotlib.pyplot as plt
from xgboost import XGBClassifier
from tqdm import tqdm
from sklearn.model_selection import LeaveOneGroupOut
from sklearn.utils.class_weight import compute_class_weight
from sklearn.metrics import (
    f1_score, accuracy_score, recall_score, precision_score,
    precision_recall_curve, confusion_matrix, roc_auc_score,
    matthews_corrcoef, roc_curve
)

warnings.filterwarnings("ignore")
```

## 0.0.1 Metric Calculation Utilities

```
[2]: def get_aupr(pre, rec):
    pr_value = 0.0
    for ii in range(len(rec[:-1])):
        x_r, x_l = rec[ii], rec[ii+1]
        y_t, y_b = pre[ii], pre[ii+1]
        tempo = abs(x_r - x_l) * (y_t + y_b) * 0.5
        pr_value += tempo
    return pr_value

def scores(y_test, y_pred, th=0.5):
    y_predlabel = [(0. if item < th else 1.) for item in y_pred]
    tn, fp, fn, tp = confusion_matrix(y_test, y_predlabel).flatten()
    SPE = tn / (tn + fp)
    MCC = matthews_corrcoef(y_test, y_predlabel)
    fpr, tpr, _ = roc_curve(y_test, y_pred)
    sen, spe, pre, f1, mcc, acc, auc, tn, fp, fn, tp = np.array([
```

```

        recall_score(y_test, y_predlabel), SPE, precision_score(y_test,
↪y_predlabel),
        f1_score(y_test, y_predlabel), MCC, accuracy_score(y_test, y_predlabel),
        roc_auc_score(y_test, y_pred), tn, fp, fn, tp
    ])
    precision, recall, _ = precision_recall_curve(y_test, y_pred)
    auapr = get_aupr(precision, recall)
    return [auapr, auc, f1, acc, sen, spe, pre, fpr, tpr, precision, recall]

```

## 1 — Utility Functions —

```

[3]: def combine_features(phage_dna, host_dna, phage_pro, host_pro):
        combined = np.concatenate([phage_dna, host_dna, phage_pro, host_pro],
↪axis=1)
        return combined

```

```

[4]: def load_feature_vector(file_path):
        return np.loadtxt(file_path)

```

```

[5]: def obtain_features(phage_list, host_list, labels, dna_base, pro_base):
        X_phage_dna, X_host_dna = [], []
        X_phage_pro, X_host_pro = [], []
        for p, h in zip(phage_list, host_list):
            X_phage_dna.append(load_feature_vector(os.path.join(dna_base,
↪'phage_dna_norm_features', f'{p}.txt')))
            X_host_dna.append(load_feature_vector(os.path.join(dna_base,
↪'host_dna_norm_features', f'{h}.txt')))
            X_phage_pro.append(load_feature_vector(os.path.join(pro_base, 'phage',
↪f'{p}.txt')))
            X_host_pro.append(load_feature_vector(os.path.join(pro_base,
↪'bacteria', f'{h}.txt')))
        return (np.array(X_phage_dna), np.array(X_host_dna),
                np.array(X_phage_pro), np.array(X_host_pro),
                np.array(labels))

```

## 2 — Load interaction matrix —

```

[6]: interaction_matrix_path = "./phage_host_interactions (1).csv"
        dna_base = './dna_features'
        pro_base = './protein_features'

        # Get list of phages and hosts based on files present
        df = pd.read_csv(interaction_matrix_path, index_col=0)

```

```

valid_phages = set([f.split('.')[0] for f in os.listdir(dna_base+"/
↳phage_dna_norm_features") if f.endswith('.txt')])
valid_hosts = set([f.split('.')[0] for f in os.listdir(pro_base+"/bacteria") if
↳f.endswith('.txt')])

df = df.loc[df.index.intersection(valid_hosts), df.columns.
↳intersection(valid_phages)]
#TODO: filter unused interaction since matrix too large

phages = df.columns.tolist()
hosts = df.index.tolist()
print(len(phages))
print(len(hosts))

```

105

200

### 3 Prepare data as list of (phage, host, label)

```

[7]: all_data = []

for p in phages:
    for h in hosts:
        label = df.loc[h, p] # Corrected access
        if pd.isna(label):
            continue # Skip missing values
        all_data.append([p, h, int(label)])

```

#### 3.0.1 Model Training and Evaluation (with Metrics)

```

[21]: results_all = []
fprs, tprs, precisions, recalls = [], [], [], []
# Your cross-validation setup:
kf = KFold(n_splits=5, random_state=1, shuffle=True)

for fold, (train_idx, val_idx) in enumerate(kf.split(all_data)):
    print(f"Fold {fold+1}")

    train_set = [all_data[i] for i in train_idx]
    val_set = [all_data[i] for i in val_idx]

    train_phages = [x[0] for x in train_set]
    train_hosts = [x[1] for x in train_set]
    train_labels = [x[2] for x in train_set]

    val_phages = [x[0] for x in val_set]

```

```

val_hosts = [x[1] for x in val_set]
val_labels = [x[2] for x in val_set]

X_phage_dna_tr, X_host_dna_tr, X_phage_pro_tr, X_host_pro_tr, y_train =   

↳ obtain_features(  

    train_phages, train_hosts, train_labels, dna_base, pro_base)  

X_phage_dna_val, X_host_dna_val, X_phage_pro_val, X_host_pro_val, y_val =   

↳ obtain_features(  

    val_phages, val_hosts, val_labels, dna_base, pro_base)  

X_train_combined = combine_features(X_phage_dna_tr, X_host_dna_tr,   

↳ X_phage_pro_tr, X_host_pro_tr)  

X_val_combined = combine_features(X_phage_dna_val, X_host_dna_val,   

↳ X_phage_pro_val, X_host_pro_val)  

imbalance = sum([1 for i in y_train if i==1]) / sum([1 for i in y_train if   

↳ i==0])  

model = XGBClassifier(  

    scale_pos_weight=1/imbalance,  

    learning_rate=0.3,  

    n_estimators=250,  

    max_depth=7,  

    use_label_encoder=False,  

    eval_metric='logloss',  

    n_jobs=8  

)  

model.fit(X_train_combined, y_train)  

y_pred_prob = model.predict_proba(X_val_combined)[:,-1]  

fold_metrics = scores(y_val, y_pred_prob)  

results_all.append(fold_metrics[:7]) # Save base metrics  

fprs.append(fold_metrics[7])  

tprs.append(fold_metrics[8])  

precisions.append(fold_metrics[9])  

recalls.append(fold_metrics[10])  

print(f"Fold {fold+1} | AUPR: {fold_metrics[0]:.4f}, AUC: {fold_metrics[1]:.  

↳ 4f}, F1: {fold_metrics[2]:.4f}, Acc: {fold_metrics[3]:.4f}")

```

Fold 1

Exception ignored on calling ctypes callback function: <bound method DataIter.\_next\_wrapper of <xgboost.data.SingleBatchInternalIter object at 0x0000025422D01480>>

Traceback (most recent call last):

File "D:\Bachelorarbeit\Prediction Notebooks\.venv2\lib\site-

```

packages\xgboost\core.py", line 640, in _next_wrapper
    return self._handle_exception(lambda: int(self.next(input_data)), 0)
File "D:\Bachelorarbeit\Prediction Notebooks\.venv2\lib\site-
packages\xgboost\core.py", line 553, in _handle_exception
    return fn()
File "D:\Bachelorarbeit\Prediction Notebooks\.venv2\lib\site-
packages\xgboost\core.py", line 640, in <lambda>
    return self._handle_exception(lambda: int(self.next(input_data)), 0)
File "D:\Bachelorarbeit\Prediction Notebooks\.venv2\lib\site-
packages\xgboost\data.py", line 1654, in next
    input_data(**self.kwargs)
File "D:\Bachelorarbeit\Prediction Notebooks\.venv2\lib\site-
packages\xgboost\core.py", line 729, in inner_f
    return func(**kwargs)
File "D:\Bachelorarbeit\Prediction Notebooks\.venv2\lib\site-
packages\xgboost\core.py", line 628, in input_data
    dispatch_proxy_set_data(self.proxy, new, cat_codes)
File "D:\Bachelorarbeit\Prediction Notebooks\.venv2\lib\site-
packages\xgboost\data.py", line 1753, in dispatch_proxy_set_data
    proxy._ref_data_from_array(data) # pylint: disable=W0212
File "D:\Bachelorarbeit\Prediction Notebooks\.venv2\lib\site-
packages\xgboost\core.py", line 1462, in _ref_data_from_array
    _check_call(_LIB.XGProxyDMatrixSetDataDense(self.handle,
array_interface(data)))
KeyboardInterrupt:

```

```

Fold 1 | AUPR: 0.2053, AUC: 0.7564, F1: 0.2737, Acc: 0.9655
Fold 2
Fold 2 | AUPR: 0.4361, AUC: 0.8058, F1: 0.4272, Acc: 0.9705
Fold 3
Fold 3 | AUPR: 0.3103, AUC: 0.8227, F1: 0.4330, Acc: 0.9725
Fold 4
Fold 4 | AUPR: 0.2924, AUC: 0.7381, F1: 0.3469, Acc: 0.9680
Fold 5
Fold 5 | AUPR: 0.3579, AUC: 0.7549, F1: 0.4000, Acc: 0.9685

```

### 3.0.2 Plotting Metrics Across Folds


```


[9]: # Plot ROC Curves
plt.figure(figsize=(10, 5))
for i in range(len(fprs)):
    plt.plot(fprs[i], tprs[i], label=f'Fold {i+1}')
plt.plot([0, 1], [0, 1], 'k--')
plt.title("ROC Curve Across Folds")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid()

```

```
plt.show()

# Plot PR Curves
plt.figure(figsize=(10, 5))
for i in range(len(precisions)):
    plt.plot(recalls[i], precisions[i], label=f'Fold {i+1}')
plt.title("Precision-Recall Curve Across Folds")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.legend()
plt.grid()
plt.show()
```

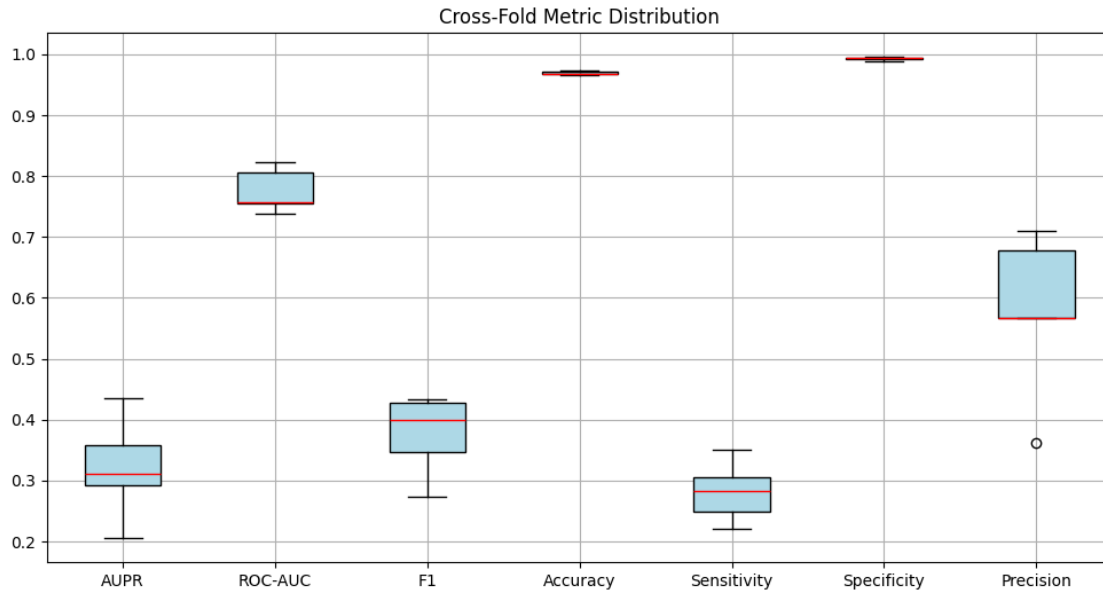




### 3.0.3 Summary Boxplot of All Metrics

```
[10]: metric_names = ["AUPR", "ROC-AUC", "F1", "Accuracy", "Sensitivity",
    ↪ "Specificity", "Precision"]
results_array = np.array(results_all)

plt.figure(figsize=(12, 6))
plt.boxplot(results_array, labels=metric_names, patch_artist=True,
            boxprops=dict(facecolor='lightblue', color='black'),
            medianprops=dict(color='red'),
            whiskerprops=dict(color='black'))
plt.title("Cross-Fold Metric Distribution")
plt.grid()
plt.show()
```



### 3.0.4 Fold-wise Metric Table

```
[11]: results_df = pd.DataFrame(results_all, columns=metric_names)
      results_df.index = [f"Fold {i+1}" for i in range(len(results_all))]
      display(results_df)

      print("Mean Metrics:")
      display(results_df.mean())
```

	AUPR	ROC-AUC	F1	Accuracy	Sensitivity	Specificity \
Fold 1	0.205312	0.756370	0.273684	0.965534	0.220339	0.988163
Fold 2	0.436058	0.805793	0.427184	0.970515	0.305556	0.995334
Fold 3	0.310329	0.822746	0.432990	0.972514	0.350000	0.991757
Fold 4	0.292393	0.738124	0.346939	0.968016	0.250000	0.993275
Fold 5	0.357870	0.754926	0.400000	0.968516	0.283784	0.994811

	Precision
Fold 1	0.361111
Fold 2	0.709677
Fold 3	0.567568
Fold 4	0.566667
Fold 5	0.677419

Mean Metrics:

AUPR	0.320392
ROC-AUC	0.775592
F1	0.376159
Accuracy	0.969019



```
Sensitivity    0.281936
Specificity    0.992668
Precision      0.576488
dtype: float64
```

### 3.1 Random Forest Model with Grid SearchCV

```
[12]: rf_results_all = []
      rf_fprs, rf_tprs, rf_precisions, rf_recalls = [], [], [], []

      rf_param_grid = {
          'n_estimators': [50, 100, 200],
          'max_depth': [None, 10],
          'min_samples_split': [2, 5],
          'min_samples_leaf': [1, 2],
          'max_features': ['auto', 'sqrt', 'log2'],
          'bootstrap': [True, False]
      }

      for fold, (train_idx, val_idx) in enumerate(kf.split(all_data)):
          print(f"RF Fold {fold+1}")

          train_set = [all_data[i] for i in train_idx]
          val_set = [all_data[i] for i in val_idx]

          train_phages = [x[0] for x in train_set]
          train_hosts = [x[1] for x in train_set]
          train_labels = [x[2] for x in train_set]

          val_phages = [x[0] for x in val_set]
          val_hosts = [x[1] for x in val_set]
          val_labels = [x[2] for x in val_set]

          X_phage_dna_tr, X_host_dna_tr, X_phage_pro_tr, X_host_pro_tr, y_train = ↵
          ↵obtain_features(
              train_phages, train_hosts, train_labels, dna_base, pro_base)
          X_phage_dna_val, X_host_dna_val, X_phage_pro_val, X_host_pro_val, y_val = ↵
          ↵obtain_features(
              val_phages, val_hosts, val_labels, dna_base, pro_base)

          X_train_combined = combine_features(X_phage_dna_tr, X_host_dna_tr, ↵
          ↵X_phage_pro_tr, X_host_pro_tr)
          X_val_combined = combine_features(X_phage_dna_val, X_host_dna_val, ↵
          ↵X_phage_pro_val, X_host_pro_val)

          class_weights = compute_class_weight(class_weight='balanced', classes=np.
          ↵array([0, 1]), y=y_train)
```

```

class_weight_dict = {0: class_weights[0], 1: class_weights[1]}

rf_base = RandomForestClassifier(class_weight=class_weight_dict,
    random_state=42, n_jobs=-1)
grid_search = GridSearchCV(estimator=rf_base, param_grid=rf_param_grid,
    scoring='f1', cv=3, n_jobs=-1, verbose=0)
grid_search.fit(X_train_combined, y_train)
best_rf = grid_search.best_estimator_

y_pred_prob = best_rf.predict_proba(X_val_combined)[: , 1]

rf_fold_metrics = scores(y_val, y_pred_prob)
rf_results_all.append(rf_fold_metrics[:7])
rf_fprs.append(rf_fold_metrics[7])
rf_tprs.append(rf_fold_metrics[8])
rf_precisions.append(rf_fold_metrics[9])
rf_recalls.append(rf_fold_metrics[10])

print(f"Fold {fold+1} | AUPR: {rf_fold_metrics[0]:.4f}, AUC:
    {rf_fold_metrics[1]:.4f}, "
        f"F1: {rf_fold_metrics[2]:.4f}, Acc: {rf_fold_metrics[3]:.4f}")

```

```

RF Fold 1
Fold 1 | AUPR: 0.1528, AUC: 0.7708, F1: 0.2553, Acc: 0.9650
RF Fold 2
Fold 2 | AUPR: 0.2967, AUC: 0.7781, F1: 0.3580, Acc: 0.9480
RF Fold 3
Fold 3 | AUPR: 0.2531, AUC: 0.8121, F1: 0.2597, Acc: 0.9715
RF Fold 4
Fold 4 | AUPR: 0.1722, AUC: 0.7137, F1: 0.2424, Acc: 0.9375
RF Fold 5
Fold 5 | AUPR: 0.3326, AUC: 0.7834, F1: 0.3371, Acc: 0.9410

```

### 3.1.1 Random Forest ROC & PR Curves

```

[13]: plt.figure(figsize=(10, 5))
    for i in range(len(rf_fprs)):
        plt.plot(rf_fprs[i], rf_tprs[i], label=f'Fold {i+1}')
    plt.plot([0, 1], [0, 1], 'k--')
    plt.title("Random Forest ROC Curve Across Folds")
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.legend()
    plt.grid()
    plt.show()



plt.figure(figsize=(10, 5))

```

```

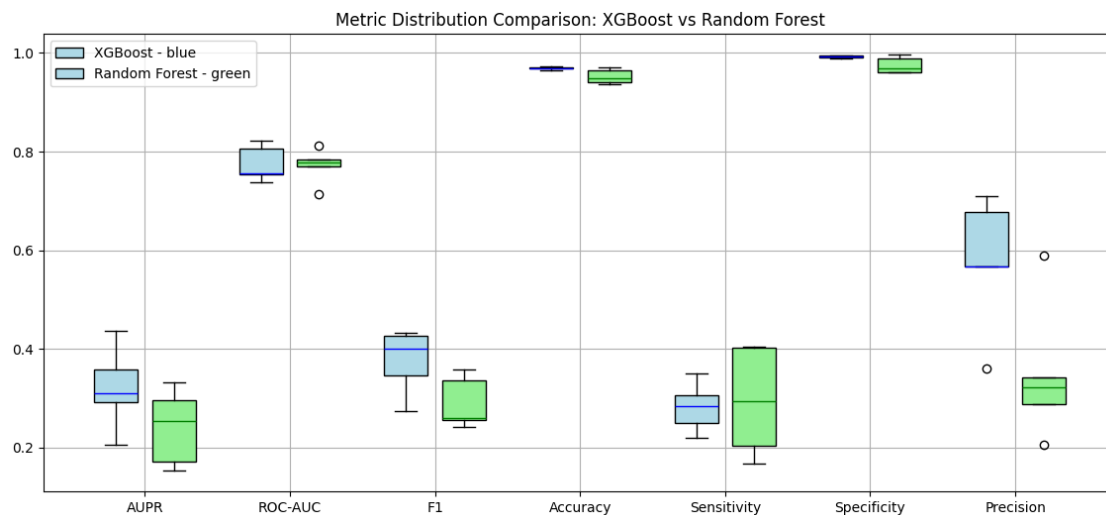
for i in range(len(rf_precisions)):
    plt.plot(rf_recalls[i], rf_precisions[i], label=f'Fold {i+1}')
plt.title("Random Forest Precision-Recall Curve Across Folds")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.legend()
plt.grid()
plt.show()

```



### 3.1.2 Metric Boxplot: XGBoost vs Random Forest

```
[14]: all_metric_names = ["AUPR", "ROC-AUC", "F1", "Accuracy", "Sensitivity",  
    ↪ "Specificity", "Precision"]  
  
plt.figure(figsize=(14, 6))  
plt.boxplot([np.array(results_all)[: ,i] for i in range(7)],  
    positions=np.arange(1, 8) - 0.2, widths=0.3, patch_artist=True,  
    boxprops=dict(facecolor='lightblue'),  
    ↪ medianprops=dict(color='blue'), labels=all_metric_names)  
  
plt.boxplot([np.array(rf_results_all)[: ,i] for i in range(7)],  
    positions=np.arange(1, 8) + 0.2, widths=0.3, patch_artist=True,  
    boxprops=dict(facecolor='lightgreen'),  
    ↪ medianprops=dict(color='green'))  
  
plt.legend(['XGBoost - blue', 'Random Forest - green'])  
plt.title("Metric Distribution Comparison: XGBoost vs Random Forest")  
plt.grid()  
plt.xticks(np.arange(1, 8), all_metric_names)  
plt.show()
```



### 3.1.3 Random Forest Fold-wise Metrics

```
[15]: rf_df = pd.DataFrame(rf_results_all, columns=all_metric_names)
      rf_df.index = [f"Fold {i+1}" for i in range(len(rf_results_all))]
      display(rf_df)

      print("Random Forest - Mean Metrics:")
      display(rf_df.mean())
```

	AUPR	ROC-AUC	F1	Accuracy	Sensitivity	Specificity \
Fold 1	0.152843	0.770798	0.255319	0.965035	0.203390	0.988163
Fold 2	0.296718	0.778087	0.358025	0.948026	0.402778	0.968377
Fold 3	0.253124	0.812142	0.259740	0.971514	0.166667	0.996394
Fold 4	0.172177	0.713741	0.242424	0.937531	0.294118	0.960166
Fold 5	0.332612	0.783447	0.337079	0.941029	0.405405	0.961598

	Precision
Fold 1	0.342857
Fold 2	0.322222
Fold 3	0.588235
Fold 4	0.206186
Fold 5	0.288462

Random Forest - Mean Metrics:

AUPR	0.241495
ROC-AUC	0.771643
F1	0.290517
Accuracy	0.952627
Sensitivity	0.294471
Specificity	0.974940
Precision	0.349592

dtype: float64

### 3.1.4 Side-by-side Mean Metrics (XGBoost vs RF)

```
[16]: xgb_mean = pd.DataFrame([np.mean(results_all, axis=0)],
      ↪ columns=all_metric_names, index=["XGBoost"])
      rf_mean = pd.DataFrame([np.mean(rf_results_all, axis=0)],
      ↪ columns=all_metric_names, index=["Random Forest"])

      display(pd.concat([xgb_mean, rf_mean]))
```

	AUPR	ROC-AUC	F1	Accuracy	Sensitivity \
XGBoost	0.320392	0.775592	0.376159	0.969019	0.281936
Random Forest	0.241495	0.771643	0.290517	0.952627	0.294471

	Specificity	Precision
XGBoost	0.992668	0.576488

Random Forest      0.974940    0.349592

## 4 XGBoost with LOGO

```
[17]: # Assuming:
# all_data = [(phage, host, label), ...]
# hosts = list of host IDs
# phages = list of phage IDs

# Create groups based on host for LOGO
groups = []
for sample in all_data:
    phage, host, label = sample
    group_id = hosts.index(host) # group by host
    groups.append(group_id)

print(f"Number of unique groups (hosts): {len(set(groups))}")
```

Number of unique groups (hosts): 200

```
[18]: logo = LeaveOneGroupOut()
cpus = 8

# For collecting predictions and labels
scores_all = []
label_list = []

# Progress bar for number of groups
pbar = tqdm(total=len(set(groups)))

for fold, (train_idx, val_idx) in enumerate(logo.split(all_data,
↪groups=groups)):
    # Prepare data
    train_set = [all_data[i] for i in train_idx]
    val_set = [all_data[i] for i in val_idx]

    train_phages = [x[0] for x in train_set]
    train_hosts = [x[1] for x in train_set]
    train_labels = [x[2] for x in train_set]

    val_phages = [x[0] for x in val_set]
    val_hosts = [x[1] for x in val_set]
    val_labels = [x[2] for x in val_set]

    # Obtain features
    X_phage_dna_tr, X_host_dna_tr, X_phage_pro_tr, X_host_pro_tr, y_train =
↪obtain_features(train_phages, train_hosts, train_labels, dna_base, pro_base)
```

```

X_phage_dna_val, X_host_dna_val, X_phage_pro_val, X_host_pro_val, y_val =
↳ obtain_features(val_phages, val_hosts, val_labels, dna_base, pro_base)

# Combine features for train and val
X_train_combined = combine_features(X_phage_dna_tr, X_host_dna_tr,
↳ X_phage_pro_tr, X_host_pro_tr)
X_val_combined = combine_features(X_phage_dna_val, X_host_dna_val,
↳ X_phage_pro_val, X_host_pro_val,)

# Class imbalance handling
pos = sum(y_train)
neg = len(y_train) - pos
scale_pos_weight = neg / pos if pos > 0 else 1

# Train
model = XGBClassifier(
    scale_pos_weight=scale_pos_weight,
    learning_rate=0.3,
    n_estimators=250,
    max_depth=7,
    use_label_encoder=False,
    eval_metric='logloss',
    n_jobs=cpus
)
model.fit(X_train_combined, y_train)

# Predict
y_pred_prob = model.predict_proba(X_val_combined)[: , 1]
scores_all.append(y_pred_prob)
label_list.append(y_val)

pbar.update(1)

pbar.close()

# Concatenate all predictions and labels
y_pred_all = np.concatenate(scores_all)
y_true_all = np.concatenate(label_list)

# Now evaluate final performance
final_metrics = scores(y_true_all, y_pred_all)

print("\nFinal LOGO metrics across all folds:")
print(f"AUPR: {final_metrics[0]:.4f}, AUC: {final_metrics[1]:.4f}, F1:
↳ {final_metrics[2]:.4f}, Accuracy: {final_metrics[3]:.4f}")

```

72%| | 144/200 [33:41<13:39, 14.63s/it]  
KeyboardInterrupt

[ ]: